



Catalyst Network

Technical White Paper

Catalyst Network - Technical White Paper

Joseph Kearney*- Atlas City

November 8, 2019

Abstract

hello

*joseph.kearney@atlascity.io

Contents

1	Distributed File System	4
2	KVM	4
2.1	From the EVM	4
2.2	To the KVM	4
3	Catalyst Consensus Mechanism	4
3.1	Notation	4
3.2	Producer Node Selection	4
3.3	Construction Phase	5
3.3.1	Local partial ledger state update generation and broadcast	6
3.3.2	Partial ledger state update collection	6
3.4	Campaigning Phase	6
3.4.1	Local candidate generation and broadcast	6
3.4.2	Candidate collection	7
3.5	Voting Phase	7
3.5.1	Ballot generation and broadcast	7
3.5.2	Ballot collection	8
3.6	Synchronisation Phase	9
3.6.1	Final ledger state update generation and broadcast	9
3.6.2	Ledger state synchronisation across the network	9

Introduction

1 Distributed File System

Catalyst integrates its own DFS based on the InterPlanetary File System (IPFS) protocol [1].

2 KVM

The Catalyst network

2.1 From the EVM

2.2 To the KVM

3 Catalyst Consensus Mechanism

On distributed networks there is no single point of trust to determine the validity of transactions, therefore concurrency must be ensured by other methods. Typically this requires a majority of the network's participants to agree on a particular update of the ledger and the account balances / holdings held within. Blockchain technologies generally employ Proof-of-Work (PoW) and occasionally Proof-of-Stake (PoS) mechanisms in order to gain consensus across a network. However, these methods are prone to increasing centralization at scale as well as in the case of PoW high energy consumption. Other networks employ a small amount of trusted nodes that ensure the validity of transactions, however this is highly centralised and almost as fallible as the single point of failure systems that DLT endeavors to avoid.

Catalyst integrates a newly designed consensus mechanism, based on Probabilistic Byzantine Fault Tolerance (PBFT). This is a collaborative rather than competitive protocol, meaning that all honest work performed by nodes on the network benefits the security of the network and that all participating nodes are rewarded equally. For each ledger cycle a random selection of worker nodes are selected, the nodes become the producers for a cycle or number of cycles. The producer nodes perform work in the form of compiling and validating transaction thereby extracting a ledger state change for that cycle.

The protocol is split into four distinct phases:

- Construction Phase - Producer nodes that have been selected create what they believe to be the correct update of the ledger. They then distribute this proposed ledger update in the form of it hash digest.
- Campaigning Phase - Producer nodes designate and declare what they believe to be the most popular ledger state update.
- Voting Phase -
- Synchronisation Phase - In this phase the producers who have computed the correct ledger update can broadcast this update the rest of the network.

3.1 Notation

3.2 Producer Node Selection

Original work for the Catalyst consensus mechanism states that the peers are selected with relation to the PID and the hash of the previous data. As the PID can be manipulated by a user and thereby weight in their favor of selection to become a producer this is not usable.

Research into a RANDAO provides a viable alternative [?][?]. Generation of a process by which each user creates their own random value and then by combining these random numbers across the network you gain distributed generated pseudo-random numbers. The larger the network the more random the number will be. The proposed process works as follows:

- Each node n in the worker pool N generates a random number r .
- To this random number the hash of the previous ledger state, D , must be added.
- n then creates a Blake-2b hash of the combined random number $H(r + D)$.
- Each n must then send their value r to the contract.
- If they do not send their r value they are not eligible to become a producer node.
- Each n in N sends their $H(r + D)$ to a hard coded smart contract. This creates the global random value R . This is done through addition of all the random numbers.
- The smart contract must determine:
 - Whether n did submit all correct information. i.e. r , and $H(r + D)$.
 - That n did in fact use the D value when generating the random number. This is done by taking the r value submitted by the user and hashing with D .
 - Ensuring n has paid a sufficient stake to take part in the selection process.
 - Validating that each worker n has only distributed one random number to the smart contract.
- Failure of any one of these points means the smart contract will not accept the submission of a random number from r and stake made will be lost.
- This global random can be used to determine the producers for the next cycle(s).
- This is done by determining the nodes that have a $H(r + D)$ closest to the R value.

This method is secure from manipulation as hashing algorithms are one way functions meaning that there is no provably efficient method to inverse a hashing function i.e. retrieving a message m from a digest $H(m)$. If a node does not input a value into the smart contract then they are not eligible for selection for becoming a producer for that cycle(s).

Addition of the D value is necessary as this will prevent a producer from using known random values to create a desired digest that gives them an advantage when being selected. The value for D must fulfill two rules, firstly it must always be the same for all nodes in the worker pool, secondly it must change with each draw of a random number when determining the random selection of producer nodes. This prevents a user creating a hash using known input and digest combinations to gain an advantage when being selected. Furthermore if the hash of the previous ledger state is used as D it ensures that a prospective producer node knows the current ledger state. If they do not then their random number will be invalid as $H(r + D_{prod}) \neq H(r + D)$.

As described in RANDAO, this can be further extended to implement staking. Nominal fees to contribute can be added in order to prevent DDOS. This is done in such a way that a nominal fee is added as to not dissuade users from legitimately wanting to become producers while dissuading malicious entities from attempting to perform a Sybil attack against the network in order to gain majority control over a producer pool for a cycle or multiple cycle. There is also the additional benefit of simplification of the overall consensus mechanism as it removes the need for a queuing mechanism as well as producing a verifiable method of keeping track of what nodes are registering to be workers for any given cycle. It also thereby in turn provides evidence to other nodes on the network who the producers for any given cycle are as the process will be verifiable.

This scheme will provide the Catalyst consensus mechanism with a verifiable, reliable and secure mechanism to generate a pool of producers. By randomly assigning which workers get to participate in the next ledger cycle fairness is ensured.

3.3 Construction Phase

The first phase of the Catalyst consensus algorithm is the Construction Phase. Within which the selected producer nodes P calculate their proposed ledger state update or their local ledger state update. This is done by aggregating and validating all transactions that have occurred during a set time period. These transactions assuming their validity are integrated into the producers local ledger state update. From which they can create a hash of the update. This hash digest represents what they believe to be the

correct update and is broadcast to the other producer nodes during for that cycle. Assuming the collision free nature of hash functions, the only mechanism for multiple producer nodes to have the same local ledger state update is by both using the same set of transactions.

3.3.1 Local partial ledger state update generation and broadcast

Each producer p in the set of producers P follows the same protocol. The construction phase begins with producer p beginning their construction phase by flushing their mempool. This mempool T is made up of n transactions t where n is the number of transactions that have been broadcast to the network and have been stored by p . These transaction are used to create p 's local ledger state update u .

The producer at this point also creates a hash trees d , this is to store the the signatures that are extracted from each transaction in T . A salt σ is created utilizing a pseudo-random number generator using the previous ledger state update $U - 1$ as its seed. p then follows the following steps:

1. Producer p verifies that each transaction in T is valid following the rules set out in [2]. From each of the n transactions in T the entries E that constitute the transaction t are extracted to form a list $E = e_1, \dots, e_m$ for m entries in t . The producer should therefore end up with n lists for E from T . each signature from the transactions are also extracted and added to d .
2. p for each E it created then creates a corresponding hash digest as:

$$L = \mathcal{H}[E \parallel \sigma]$$

Each pair (E, O) is added to a list L .

3. p then sorts list L into lexicographical order according the hash values O .
4. The producer p then extracts the transaction fee value from each transaction in T to create v which is the total sum of all transaction fees.
5. The local ledger state update u for producer p can then be calculated. Firstly the list L is concatenated (denoted by \parallel) with the hashtree d and a hash digest is created as:

$$u = \mathcal{H}(L \parallel d)$$

u is then concatenated with p 's unique peer identifier Id to create:

$$h = u \parallel Id$$

6. h is then broadcast to the other producer nodes on the network.

3.3.2 Partial ledger state update collection

Producer p also collects other producers partial ledger update values. At most they will collect $P - 1$ values. Optimally every producer in P will receive the same set of transactions therefore for every p in P will have the same partial ledger update u . However this is unlikely due to all transactions not being received by a small group of nodes. Equally they may not hold G where $G = h_1, \dots, h_P$, meaning they may not receive a proposed update from all candidates.

3.4 Campaigning Phase

The second phase of the consensus mechanism is where producer p designated as candidate for what it calculates to be the most popular ledger state update.

3.4.1 Local candidate generation and broadcast

Beginning this phase producer p has a set of partial ledger updates G that it has received from other producer nodes. Each h within G contains a producers hash of the proposed update (u and a peer identifier (Id)). The most popular u value can be found, this gives us u^{maj} from there the subset G_{maj} can be created, which is the amount of votes for the most popular update. Two thresholds must be considered first is G_{min} this is the minimum amount of updates it has received from other producers

in order to generate a valid candidate. The second is G_{thresh} this is the threshold value for which a minimum number of votes must be in favor of G_{maj} which the most popular vote found within G . So in order to proceed with declaring a candidate $G > G_{min}$ and $G_{maj} > G_{thresh}$.

If the thresholds are met the following can take place:

1. p creates a list $\mathcal{L}(prod)$. To this list p appends the identifier of any producer that correctly sent the u value that equals u^{maj} . If p 's u value is also the same as u^{maj} then they should append their own Id .
2. Producer p then creates their candidate for the ledger update c which is calculated as $c = u^{maj} || \#(\mathcal{L}(prod)) || Id$
3. Producer p will then broadcast their preferred update c to the other producers.

3.4.2 Candidate collection

p during this phase will be collecting the c values from other producers. At the end of this phase of the cycle p will hold a set of C candidates.

3.5 Voting Phase

The third phase of the ledger cycle is the Voting Phase within which a producer p from the C candidates it has received decides on what it believes should be the global ledger state update i.e. the update that should be applied to the ledger for that cycle.

3.5.1 Ballot generation and broadcast

1. p verifies that the same first hash value u^{maj} is embedded in a majority of producer candidates.
2. p is required at this point to have created the majority ledger update. As only with this partial update can they participate effectively. Therefore their $u == u^{maj}$

If each producer collects the first hash value generated by every producer, any two producers would build the same list of identifiers $\mathcal{L}(prod)$. However even in highly efficient distributed network it is unlikely that all producers will retrieve all information. Therefore it must be assumed that the list of identifiers held by p is incomplete, p however must ensure that they hold enough information to confidently issue a vote on the update to the ledger. The identifier of a producer is appended to the hash they distribute because:

- To verify that p is a producer node and that they infact were one of the winners of the random draw described in ??.
- To evaluate and track the quality of the work performed by p .
- To ensure that the correct producers withing P are rewarded for their work maintaining the ledger.

New transaction entries (coinbase entries) are created using the list's $\mathcal{L}(prod)$ produced by the producers p . This allows reward and fees to be paid to the correct users. These new transaction entries will be appended to the final ledger state update for that cycle. Therefore a complete ledger state update should consist of:

- The list of transaction entries integrated into the new ledger state all E lists.
- The transaction signatures from T held in d .
- Compensation entries rewarding the producers from $\mathcal{L}(prod)$.

The voting process thereby must confirm the correct list of identifiers involved in producing the correct ledger state updates. The final list for $\mathcal{L}(prod)$ is generated by merging the producers lists together. In order for a producer p to have their Id added to the final $\mathcal{L}(prod)$ the must appear in at least 50% of the $\mathcal{L}(prod)$ lists distributed by P . This ensures that no producer can just append their own Id to their update whether that update was correct or not and gain a reward for it.

As explained below the final list $\mathcal{L}_n(prod)$ is obtained by merging the partial lists included in the producers' candidate. A producer P_j could have produced a first hash value h_{Δ_j} different to $h_{\Delta_j}^{maj}$ yet added

his identifier to $\mathcal{L}_j(prod)$ when building its candidate c_j in the attempt to collect some token reward. In such scenario Id_j would be an element of the list included in c_j (or any other producer node controlled by P_j), but it wouldn't be included in any other list $\{\mathcal{L}_k(prod)\} \forall k \in P/j$. To prevent such malicious behaviour, a rule imposes that P_j only appends to the final list $\mathcal{L}_n(prod)$ the identifier of a producer included in the list $\mathcal{L}_k(prod)$ of a candidate c_k satisfying $h_{\Delta_k}^{maj} = h^{maj}$ if and only if that identifier is included in at least $P/2$ lists $\{\mathcal{L}_k(prod)\}_{k=1,\dots,V_j}$ associated to a candidate c_k satisfying $h_{\Delta_k}^{maj} = h^{maj}$. Only a producer controlling half or more of the producer nodes would succeed in including its identifier into the final list $\mathcal{L}_n(prod)$.

Although this eliminates the risk of unethical behavior from the producer, this also means that there would be little incentive for a producer to broadcast its vote if its identifier was not included in $\mathcal{L}_n(prod)$. However, the probability that a producer compiles the correct final list $\mathcal{L}_n(prod)$ strongly depends on the number of votes collected. The more votes collected by a producer, the greater the probability that said producer will compile the complete final list. Although a producer may not have produced the correct partial ledger state update, participating in the voting process is, therefore, an important contribution to the overall consensus protocol and should entitle the producer nodes to some reward. To that end a producer P_j can use the identifier of other producers included in their vote and create a second list $\mathcal{L}_j(vote)$ to account for their participation in the voting process.

P_j follows a series of step for a period of time Δt_{v0} ($\Delta t_{v0} < \Delta t_v$):

1. P_j creates a new list $\mathcal{L}_j(vote)$ and appends to said list the identifier of any producer P_k who forwarded a candidate c_k satisfying $h_{\Delta_k}^{maj} = h^{maj}$.
2. P_j creates the final list $\mathcal{L}_n(prod)$ and appends to said list the identifier of a producer included in the list $\mathcal{L}_k(prod)$ of a candidate c_k satisfying $h_{\Delta_k}^{maj} = h^{maj}$ if and only if that identifier is included in at least $P/2$ lists $\{\mathcal{L}_k(prod)\}_{k=1,\dots,V_j}$ associated to a candidate c_k satisfying $h_{\Delta_k}^{maj} = h^{maj}$.
3. P_j then creates a list L_{CE} of compensation entries for each producer whose identifier is included in $\mathcal{L}_n(prod)$. Each producer receives x_h tokens. Assume that $C_n \leq P$ identifiers are included in $\mathcal{L}_n(prod)$ and X is the total number of tokens injected per cycle for the pool of P producers. The quantity x_h is defined such that $C_n x_h = f_{prod} X + x_f$ where x_f represents the total number of fees collected from the m_{n-1} transactions and f_{prod} represents the fraction of new tokens injected per cycle and distributed to the producers who built the correct ledger state update. The remaining $(1 - f_{prod})X$ tokens are distributed to other contributing nodes in the network. A part of this remainder goes to the producers who voted correctly on the previous ledger cycle update. Let $\mathcal{L}_{n-1}(vote)$ be the list of the identifiers of producers who voted correctly on the previous ledger cycle update \mathcal{C}_{n-1} . We later demonstrate how such a list is derived during a ledger cycle. For now, let's assume that L_{CE} includes compensation entries for producers involved the production of the ledger state update for this ledger cycle \mathcal{C}_n and the producers involved in the voting process of the preceding cycle \mathcal{C}_{n-1} .
4. P_j then creates the candidate ledger state update for \mathcal{C}_n including the reward allocated to the producers for their contribution:

$$LSU_j = L_E^f \parallel d_n \parallel L_{CE}$$

P_j then computes its vote (or *producer vote*):

$$v_j = \mathcal{H}(LSU_j) \parallel \#(\mathcal{L}_j(vote)) \parallel Id_j \quad (1)$$

which includes the hash of the candidate ledger state update (or *second hash value*) and a partial list of identifiers of producers who designated the correct candidate partial ledger state update corresponding to h^{maj} .

5. P_j then forwards v_j to the other producers and collects the producer votes issued by its peers. Figure ?? illustrates the different steps followed by P_j during the voting phase.

3.5.2 Ballot collection

During the voting phase, the producer P_j collects the producer votes broadcast by its peers. At the end of the voting phase ($t = t_v + \Delta t_v$), the producer P_j holds U_j producer votes in its cache with $U_j \leq C_n$ where $C_n \leq P$ is the actual total number of producers who correctly computed h^{maj} .

3.6 Synchronisation Phase

3.6.1 Final ledger state update generation and broadcast

The last phase (*a.k.a* synchronisation phase) of a ledger cycle starts at $t = t_s$, with $t_s = t_{n,0} + \Delta t_p + \Delta t_c + \Delta t_v$, and lasts for a period of time Δt_s , therefore ending at $t_s + \Delta t_s = t_{n,0} + \Delta t_{cycle}$.

During a period of time $\Delta t_{s0} < \Delta t_s$, P_j executes the following steps:

1. P_j defines the ledger state update ΔL_n for the cycle \mathcal{C}_n as:
 $\mathcal{H}(\Delta L_n) = \max[\text{unique}(\mathcal{H}(LSU_k)) \mid \forall k \in \{U_j\}]$ and the associated number of votes collected:
 $U^{maj} = \text{count}[(\mathcal{H}(LSU_k) = \mathcal{H}(\Delta L_n)) \mid \forall k \in \{U_j\}]$ and verifies that $U^{maj} > U_{threshold}$.
2. P_j creates a new list $\mathcal{L}_n(\text{vote})$ and append to $\mathcal{L}_n(\text{vote})$ the identifier of a producer included in the list $\mathcal{L}_k(\text{vote})$ of a vote v_k satisfying $\mathcal{H}(LSU_k) = \mathcal{H}(\Delta L_n)$ if and only if the identifier is included in at least $C_n/2$ lists $\{\mathcal{L}_k(\text{vote})\}$ associated to a producer vote v_k satisfying $\mathcal{H}(LSU_k) = \mathcal{H}(\Delta L_n)$. Note that C_n can be easily computed as it corresponds to the number of producer identifiers who correctly computed the ledger state update and are therefore included in $\mathcal{L}_n(\text{prod})$.
3. If P_j generated the correct ledger state update ΔL_n , it can write it to DFS which will return it with a content-based address \mathcal{A}_n .
4. A Producer P_j then creates the following output quantity (or *producer output*):

$$\boxed{o_j = \mathcal{A}_n \parallel \#(\mathcal{L}_n(\text{vote})) \parallel \text{Id}_j} \quad (2)$$

The producer then broadcasts o_j to the network.

3.6.2 Ledger state synchronisation across the network

During the time period $[t_s, t_s + \Delta t_{cycle}]$, user nodes collect $\{o_k\}_{\forall k \in P}$ producer outputs broadcast by the producers. By extracting the identifier Id_k embedded in any collected output o_k , a user node can easily compile a list of producer identifiers having broadcast the same second hash value $\mathcal{H}(\Delta L_n)$ (concatenated with the same list $\mathcal{L}_n(\text{vote})$). Upon receiving $x > P/2$ identical addresses $\{\mathcal{A}_k = \mathcal{A}_n\}_{k \in x}$, the user nodes can read the common address content (ΔL_n) from DFS. Using ΔL_n a user node can safely synchronise their local copy of the ledger and write it to their DFS if not already done. The balance of accounts stored on the ledger are updated and the producers effectively collect their rewards.

Worker nodes also store the list $\mathcal{L}_n(\text{vote})$ embedded in each o_k output. If selected to be a producer for the next cycle \mathcal{C}_{n+1} , a worker can use it to generate the reward allocated to the producers who correctly voted for the accurate ledger state update during the ledger cycle \mathcal{C}_n .

Figure ?? summarises the different phases of the ledger cycle.

The various parameters and thresholds mentioned in this chapter and their impact on the levels of security and confidence in the successful production of a ledger state update are discussed in section ??.

References

- [1] J. Benet, "Ipfes-content addressed, versioned, p2p file system," *arXiv preprint arXiv:1407.3561*, 2014.
- [2] catalyst network, "Transactionvalidator.cs." <https://github.com/catalyst-network/Catalyst.Framework/blob/f23089d75ce876008ae1400f8044f538e07a3100/src/Catalyst.Core.Lib/Validators/TransactionValidator.cs>, 07/11/2019.